

Computer Architecture 2025

Fall Lab 3

1. Problem Description

- In this lab, you're going to **implement a cache** for a single-cycle CPU.
 - CPU features:
 - 32 registers
 - Input: 32-bit binary instructions
 - Area Limitation (**Do not be concerned with this constraint.**)
 - CPU: 12000 mm²(In HW1, the maximum difference in area between the TA's result and any student submission is 1000. TA implemented Lab 3's CPU with 9800 mm²)
 - Cache: 20000 mm²
-

1.1 Data Path

- Figure 1 describes the CPU data path.
 - **Provided by TAs: Registers and Memory (written in TB)**
 - The Data Memory has read/write latency; you are asked to implement a cache to accelerate it.
- The single-cycle CPU architecture builds upon the HW1 implementation. However, you are required to support only the functionality explicitly covered in the provided test cases.
 - **R-type:** ADD, SUB, AND, XOR, MUL
 - **I-type:** ADDI, SLLI, SLTI, SRAI, LW, JALR
 - **S-type:** SW

- **B-type:** BEQ, BNE, BLT, BGE
- **U-type:** AUIPC
- **J-type:** JAL
- **System:** ECALL (Triggers `o_finish, 00000073`)
- Data memory
 - Data memory has a read/write latency of 10 cycles.
 - Ensure the **Read/Write (R/W)** control signals for the data memory are held stable until the data is successfully returned.
 - A block in data memory is 128 bits, containing four 32-bit words (or "distinct data") per request.
 - The last 4 bits of the data memory address are structured as: **2 bits for the block offset and 2 bits for the byte offset. (Note: You may set all 4 bits to 0 in `cache.v` when addressing the main memory.)**

```
assign o_mem_addr = {tag, index, 4'b0000} + i_offset;
```

- Cache
 - Used to accelerate **Data Memory**
 - Check the TA's implementation in the reference.
 - Be careful when handling the processor address. Before performing any address operations on `i_proc_addr`, **subtract** `i_offset`. Before outputting to the data memory (`o_DMEM_addr`), **add** the offset back.

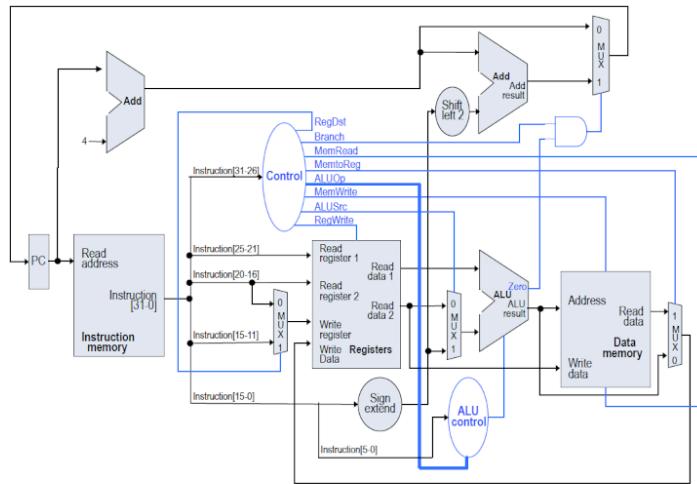


Figure 1: Referenced Single Cycle CPU

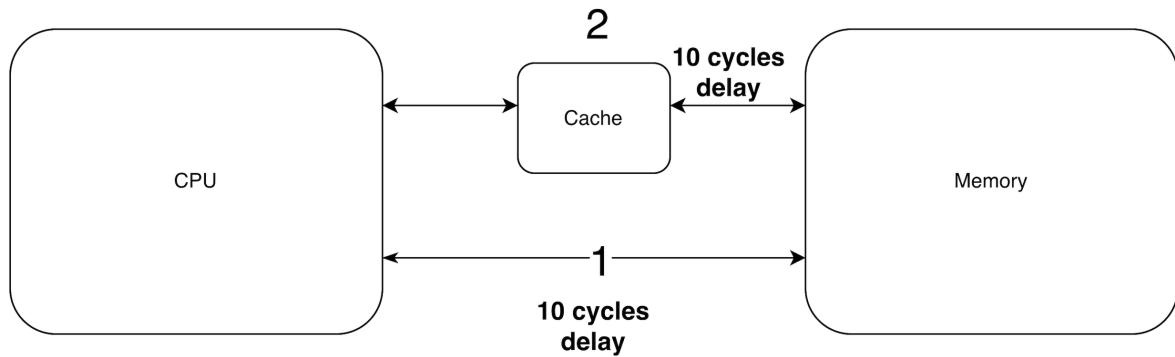


Figure 2: Memory Hierarchy and Delay

1.2 Cache

- You can choose any cache implementation
 - Write through, Write back
 - Write around

Modify the value `o_cache_available` to 1 in your implementation once cache is implemented.

```

assign o_cache_available = 1; // change this value to 1 if the cache is implemented

//-----//
|    // default connection          //
// assign o_mem_cen = i_proc_cen;      //
// assign o_mem_wen = i_proc_wen;      //
// assign o_mem_addr = i_proc_addr;    //
// assign o_mem_wdata = i_proc_wdata;   //
// assign o_proc_rdata = i_mem_rdata[0+:BIT_W];//
// assign o_proc_stall = i_mem_stall;   //
//-----//

```

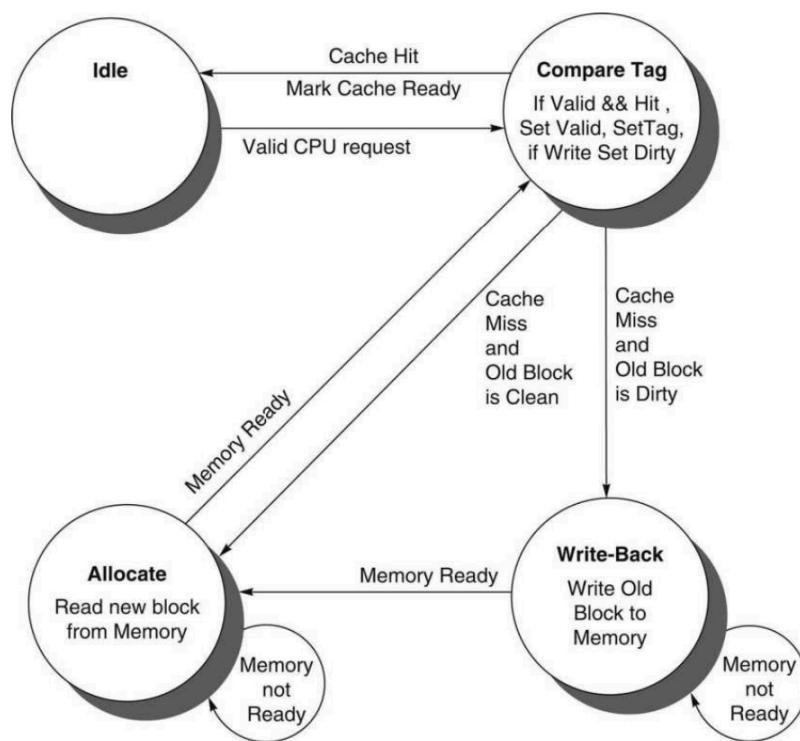


FIGURE 5.39 Four states of the simple controller.

Easy FSM for the cache implementation

1.3 Input / Output Data

- Provided: `testbench.v`, `mem_I.dat`, `mem_D.dat`, `golden.dat`, `mem_I_listing` (In I3, I4)
- All test cases should finish in 10000 cycles.

- You may change the test pattern by executing Python in the `./code/Gen_pattern`
 - I1: factorial
 - I2: arithmetic operations
 - I3: insertion sort
 - I4: bubble sort
 - All assembly code finishes with **ECALL** (00000073)
-

1.4 Required Modules

- **CPU (single cycle)**
- Cache
- **Others**
 - You may add extra helper modules if needed.

1.5 Synthesizability and Area

- All RTL code must be synthesizable (no unintended latches).
 - The TA will check your design using **yosys** (`synth.log`). You do not need to modify any yosys scripts. The results can be found in Section 10.6 of `./log/synth.log`, including the estimated area of your module.
 - Run `parse.py` to get area and latch results for CPU and cache, paste the results, and fill in the numbers in the report
-

```
(base) weiber@pasLab60:~/CA/CA-HW5/release$ =====
CPU Synthesis Log Analysis
=====
== Module Areas ==
ALU: 4400.704
ALUcontrol: 95.760
CPU: 5313.616
Control: 25.536
=====
Total area: 9835.616

== Latch Statistics (PROC_DLATCH pass) ==
No latch inferred count : 16
Latch inferred count     : 0

=====
Cache Synthesis Log Analysis
=====
== Module Areas ==
Cache: 11356.604
=====
Total area: 11356.604

== Latch Statistics (PROC_DLATCH pass) ==
No latch inferred count : 65
Latch inferred count     : 0
```

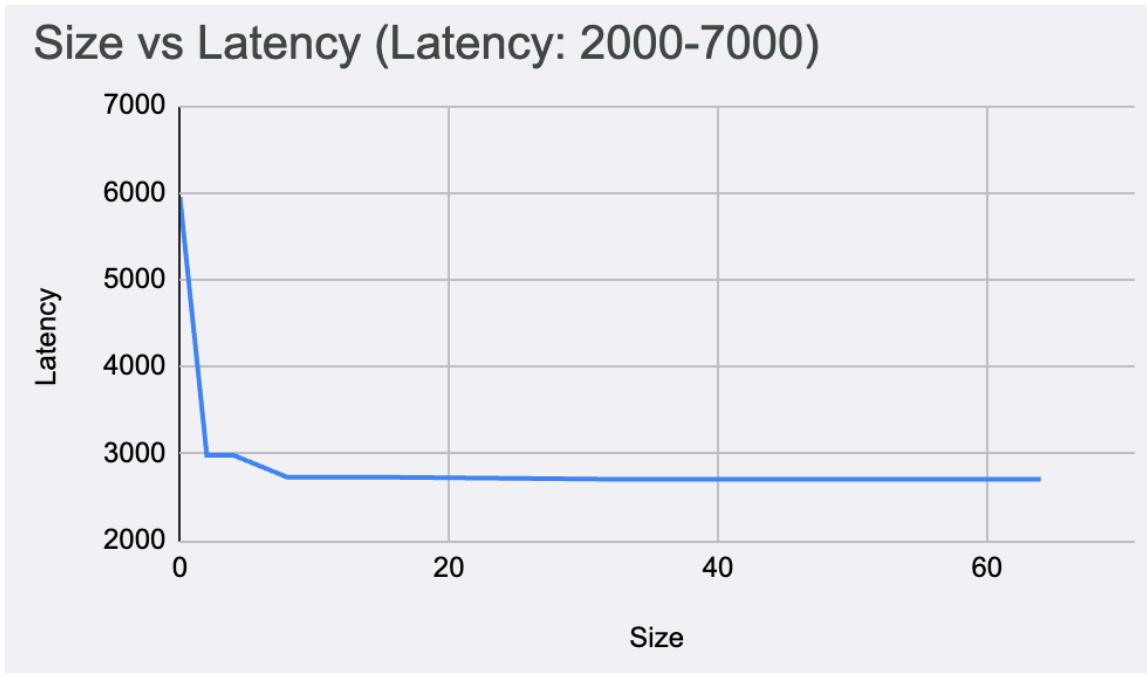
1.6 Testcases

- TAs will test your modules with public and hidden patterns.
 - Create your own test pattern through the provided Python!
-

2.1 Report

- **Explain your CPU and Cache.**
- The report also includes:
 - Synthesizability / latch check screenshot.
 - Development environment (OS, compiler, IDE).
- **Pareto figure** for cache size versus execution cycles (Draw the figure with the **I4 test case**).
 - Clearly specify the **number of different cache sizes** tested.
 - You must evaluate **at least four distinct cache sizes**.
 - Paste the “pass for simulation results” for each cache size. (see below)
 - You don’t need to use a public testcase if the figure is not clear enough.

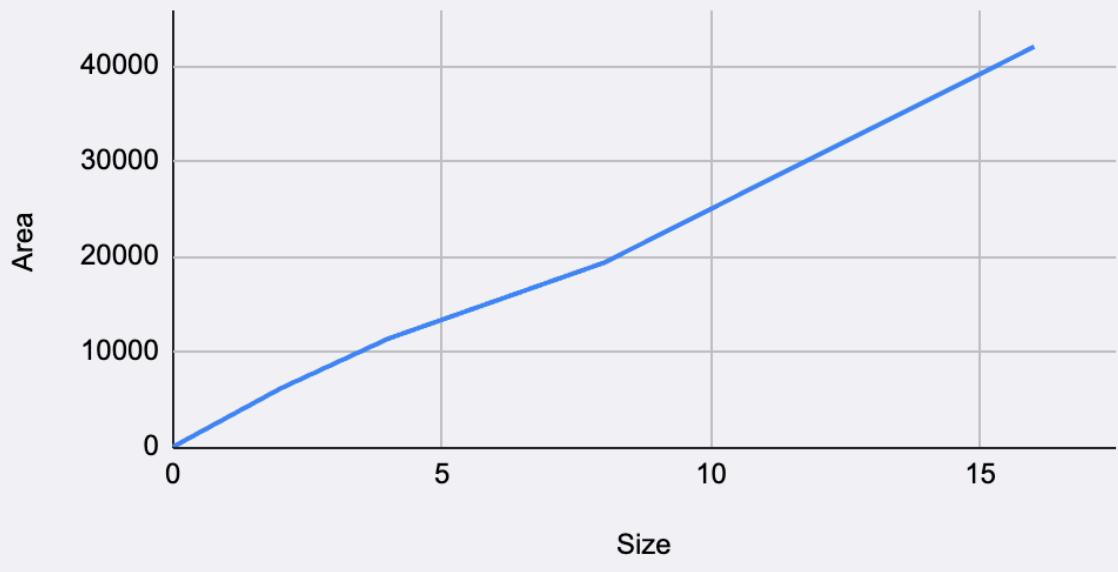
- TA reserves the right to assess and grade based on the quality of the figures, but you can add a text description for clarity.



```
-----  
START!!! I4 Simulation Start .....  
-----  
=====  
  
Success!  
The test result is .....PASS :)  
Total execution cycle : 995  
=====
```

- Figure for cache size versus cache area
 - TA reserves the right to assess and grade based on the quality of the figures.

Size vs Area



3.1 Development Environment

Requirement

- Docker is required.

Directory Structure (after unzip)

```
Lab3/
├── code
│   ├── Gen_Pattern
│   │   ├── I1_fact_gen.py
│   │   ├── I2_hw1_gen.py
│   │   ├── I3_insertion_sort_gen.py
│   │   └── I4_bubble_sort_gen.py
│   └── Pattern
│       ├── I1
│       │   ├── golden.dat
│       │   ├── mem_D.dat
│       │   └── mem_I.dat
│       └── I2
```

```
    └── golden.dat
    └── mem_D.dat
    └── mem_I.dat
    └── I3
        └── golden.dat
        └── mem_D.dat
        └── mem_I.dat
        └── mem_I_listing.txt
    └── I4
        └── golden.dat
        └── mem_D.dat
        └── mem_I.dat
        └── mem_I_listing.txt
    └── src
        └── cache.v
        └── CHIP.v
    └── supplied
        └── memory.v
        └── Regfile.v
    └── tb
        └── tb.v
    └── docker-compose.yml
    └── dockerfile
    └── judge.yaml
    └── log
        └── cache_syn.log
        └── cpu_syn.log
        └── output_1.txt
        └── output_2.txt
        └── output_3.txt
        └── output_4.txt
    └── Makefile
    └── parse.py # check for latch and area
    └── run_all.sh # use for simulation
```

 You should modify files in `code/src/`

3.2 Run with Docker

After implementation, execute:

```
sudo make run
```

- Results (logs, text, waveforms) will be generated under [log](#).

3.3 Run without Docker (alternative)

If you don't have Docker, run on **Ubuntu 22.04** with **iverilog 11.0-1.1 & Yosys & nangate45**

```
# Simulation
./run_all.sh

# 合成
yosys -l <log::path>/cpu_syn.log -p "
    read_liberty -lib <lib::path>/NangateOpenCellLibrary_typical.lib;
    read_verilog <source::path>/*.v;
    hierarchy -top CPU;
    proc; opt_clean;
    fsm; opt_clean;
    techmap; opt_clean;
    flatten CPU;
    dfflibmap -liberty <lib::path>/NangateOpenCellLibrary_typical.lib;
    abc -liberty <lib::path>/NangateOpenCellLibrary_typical.lib;
    stat -liberty <lib::path>/NangateOpenCellLibrary_typical.lib;
"
yosys -l <log::path>/cache_syn.log -p "
    read_liberty -lib <lib::path>/NangateOpenCellLibrary_typical.lib;
    read_verilog <source::path>/cache.v;
    hierarchy -top Cache;
    proc; opt_clean;
    fsm; opt_clean;
    techmap; opt_clean;
```

```
flatten Cache;  
dfflibmap -liberty <lib::path>/NangateOpenCellLibrary_typical.lib;  
abc -liberty <lib::path>/NangateOpenCellLibrary_typical.lib;  
stat -liberty <lib::path>/NangateOpenCellLibrary_typical.lib;  
"  
# area results  
python parse.py
```

⚠ Note: You will get **0 points** if your code cannot run correctly on the official environment

4.1 Submission Rules

- Directory structure after unzipping:

```
studentID_lab3/  
    └── src/ (all Verilog codes you wrote)  
    └── studentID_lab2_report.pdf
```

- Do **not** include: tb.v, memory.v, Regfile.v
- File name: **studentID_lab3.zip**.
- Submit via NTU COOL.
- Deadline: 2025/12/9 (Tue.) 23:59.

5.1 Evaluation Criteria

- **Pass Testcase (40%)**
 - Public testcases: 20% ($4 \times 5\%$)
 - Hidden testcases: 20% ($4 \times 5\%$)
- **Performance (30%)**
 - Baseline (20%) : Total execution cycles in public testcases (I1~I4) less than 5000 (implement cache)

- Ranking (10%): Shortest total execution cycles
 - You get 0 in ranking once:
 - late submission
 - violate area limitations
 - have latch violations
 - the design is unsynthesizable
 - compile error
 - Score with linear ranking
- **Report (30%)**
 - Description (10%)
 - Describe your CPU and Cache design
 - If you use LLM, clearly state which parts you use.
 - We take Plagiarism seriously.
 - Pareto and area diagram: 10% (5% for each)
 - Demo 10 %
- **Bonus (20%)**
 - TA strongly recommends a novel or customized design of the cache. If you do exceptionally well on this, no worry about the ranking score.
 - If you implement the design from paper or some complicated cache, feel free to write in the report!
 - Judged by the TA, however, implementing the cache mentioned in the class won't get any bonuses.
- **Other penalties:**
 - Compilation error (naming, or minor errors in 3 lines): -5%
 - Submitting unnecessary files (tb or supplied): -5%

- Wrong directory format: -5%.
 - Plagiarism: get 0 in this homework, and TA will inform teacher
 - Late submission: -10% per day, at most a week
 - Major mistakes causing compilation error → programming part 0. (Judged by TA)
-

Reference

This is the TA's implementation: Write through, Write around, and Direct-map cache.

```
localparam CACHE_SIZE = 4; // num of blocks
localparam INDEX_BITS = $clog2(CACHE_SIZE);
localparam TAG_BITS    = ADDR_W - INDEX_BITS - 4; // 32-2-4=26bits
// (4: 2 bits for choosing data among byte data, 2 bits for block offset)
assign addr = i_proc_addr-i_offset;
assign tag   = addr[ADDR_W-1-:TAG_BITS];
assign index = addr[(ADDR_W-1)-TAG_BITS-:INDEX_BITS];
assign block = addr[3:2];

assign o_mem_addr  = {tag, index, 4'b0000} + i_offset;
// 2 bits for choosing data among byte data and 2 bits for block offset

reg [TAG_BITS + 4*BIT_W + 1:0] cache_r [0:CACHE_SIZE-1];
reg [TAG_BITS + 4*BIT_W + 1:0] cache_w [0:CACHE_SIZE-1];
```

Size	Latency	Area	I3
64	2706		
32	2706		
16	2728	41993	416
8	2728	19311	416
4	2981	11356	416
2	2981	6133	438
0	5951	0	674

It is not the golden testcase.

	W/o cache	W cache
I1	806	692
I2	858	744
I3	2008	1098
I4	2051	995
Total	5723	3529

area:11356

It is the golden testcase (4 entry cache).

Contact

Questions: email eclab.ca.ta@gmail.com